Pension Governance Council – One-Page Overview

Purpose

Translate opaque AI systems used in pensions and asset management—LLMs, risk engines, optimizers, robo-advisers, valuation models—into board-grade accountability. This GPT helps fiduciaries interrogate models, document oversight, and keep decision rights with the board.

Who It's For

- Pension trustees, CIOs/CROs, chief actuaries, compliance and audit leads.
- Corporate boards and governance committees overseeing benefit plans or investment entities.
- Regulators/supervisors assessing model risk practices.
- Asset owners/LPs diligencing GP analytics and valuation tooling.

How to Use It (quick path)

1. **Describe the model or process.** Paste a short summary of what it does, the decisions it influences, and where it runs (in-house/vendor).
2. **Run the modules:**
    - **Model Inventory & Duty Trace** → creates a Model Accountability Sheet (owner, purpose, materiality, sign-off, regulatory exposure).
    - **Explainability & Bias Interrogator** → generates board questions on inputs, reproducibility, failure modes, and human-override points.
    - **AI Governance Playbook** → issues a lightweight policy: approval gates, validation cadence, documentation standards, escalation triggers, trustee literacy.
    - **Regulatory Readiness Scan** → highlights scrutiny points, missing artefacts, and disclosures likely required to members/regulators.
    - **Red-Team / Failure-Mode Generator** → proposes 2–3 plausible breakdowns (technical, governance, legal) with mitigations.
3. **Export the output** into board packs, risk registers, or audit trails. Treat it as working papers to be validated by your internal experts.

What It Covers (typical models)

- Longevity/mortality tables and demographic risk engines.
- LDI/asset-allocation optimizers; climate-scenario simulators.

- Employer covenant scoring; fraud/anomaly detection; robo-advice flows.
- Private-market valuation benchmarks and NLP contract review tools.

Limitations & Guardrails

- **Not legal, actuarial, or investment advice.** Outputs are governance aids and should be independently validated.
- **No automatic execution.** Any auto-trading or auto-allocation without human sign-off is flagged as a governance breach.
- **Data constraints.** Works from what you provide; it does not access your systems or proprietary vendor code. Avoid pasting sensitive data you are not permitted to share.
- **Model risk honesty.** This GPT never declares a model "safe"—only governable under stated assumptions and documented controls.
- **Fiduciary reality check.** Trustees can outsource computation, not judgment; responsibility is non-delegable.

Creator

Tanya Matanda — governance practitioner and author of the *Autumn Governance Series* on pension risk, AI assurance, and fiduciary oversight (e.g., *Technology, Data, and AI: Building Smarter Pensions*; *Fiduciary Capital: The Canadian Advantage*; *Pension Systems at the Edge*). She builds practical frameworks that help boards audit algorithms and defend decisions.

Link to the GPT

Insert your share link here: Pension Governance Council

> If sharing internally, also direct users to the GPTs directory in ChatGPT and search **"Pension Governance Council."**

Tagline: AI oversight for boards, not engineers.